

Logical Graphs

Control Flow Operations in TensorFlow



Hello!

I AM SAM ABRAHAMS

Co-author of *TensorFlow for Machine Intelligence*

Teach “Deep Learning with TensorFlow” at Metis

Long time TensorFlow contributor


The “TensorFlow on Raspberry Pi” guy (who isn’t Pete Warden)

Copyright 2017 Sam Abrahams. All Rights Reserved.


Slides and code:

github.com/samjabrahams/talks/tree/master/tensorflow/control_flow


Control flow: an example




We want to run either B or C, based on A's value



How might we do this?




Naively: use Python if/else and multiple runs,




```
a = sess.run(a_op)
feed = {a_op: a}
if a > 0:
    sess.run(b_op, feed)
else:
    sess.run(c_op, feed)
```

But this is **awkward**




```
a = sess.run(a_op)
feed = {a_op: a}
if a > 0:
    sess.run(b_op, feed)
else:
    sess.run(c_op, feed)
```


We fetch a value only to feed it back in




```
a = sess.run(a_op)
feed = {a_op: a}
if a > 0:
    sess.run(b_op, feed)
else:
    sess.run(c_op, feed)
```

Python logic **isn't represented** in the graph




```
a = sess.run(a_op)
feed = {a_op: a}
if a > 0:
    sess.run(b_op, feed)
else:
    sess.run(c_op, feed)
```

Also: one `sess.run` should represent an **entire run**



```
a = sess.run(a_op)
feed = {a_op: a}
if a > 0:
    sess.run(b_op, feed)
else:
    sess.run(c_op, feed)
```

What we want: **native logic gate**



Obvious follow up

TensorFlow has several operations for
native control flow

114

■ Why care about native control flow?

1. Efficiency
2. Flexibility
3. Compatibility

■ Efficiency

- Passing data to/from the Python layer is **slow**
- Want to **run graph end-to-end** as much as we can
- Takes advantage of pipelining, such as **queues**

■ Flexibility

- Empower static graphs with **dynamic components**
- Model logic kept in one place → **better decoupling**
- Graph can change without affecting training loop

Compatibility

- Debug and inspect with TensorBoard
- Seamlessly deploy with TensorFlow Serving
- Auto-differentiation, queues, pipelining

Note:

I'm bad with colors

Color change in

3

Color change in

2

Color change in


1




CONTROL DEPENDENCIES

Dependencies: quick recap


Here's a graph!




Nodes (operations)




Edges: (tensors)




Dependencies: nodes required to compute another node



The dependencies of the last node




Dependency of an earlier node




Dependencies and execution order

- TF keeps track of every operation's dependencies
- Uses them to **schedule computation**
 - An op is eligible to run once its dependencies have finished
- Two eligible ops can execute in **any order**


Back to our graph




These two nodes have **no dependencies**




This means they can run in **any order**




Scenario: race-condition



One operation **changes** a variable



The other **reads** from that variable




Code might look like this...


```
var = tf.Variable(...)  
top = var * 2  
bot = var.assign_add(2)  
out = top + bot
```

Note: `assign_add()` returns value of Variable after being adjusted.

Currently: execution order is **non-deterministic**



Might lead to unexpected behavior!




How do we control this?


Dependency management

- TensorFlow automatically determines dependencies
 - Basically, all of an Op's inputs
- User can define **additional dependencies**
 - Forces specified operations to complete first


We can control the order depending on needs



If we want the variable to change and then read from it, make the top depend on the bottom



If we want to read the variable before it changes,
make the bottom depend on the top



We do this with `tf.control_dependencies`

```
# Before the changes  
var = tf.Variable(...)  
top = var * 2  
bot = var.assign_add(2)  
out = top + bot
```

We do this with `tf.control_dependencies`

```
# Force bot to wait for top
var = tf.Variable(...)
top = var * 2
with tf.control_dependencies([top]):
    bot = var.assign_add(2)
out = top + bot
```


■ `tf.control_dependencies(control_inputs)`

1. Put **list of desired dependencies** as `control_inputs`
2. Ops defined in the `with` block **gain those dependencies**


Use cases

- Enforcing **execution order**
 - As shown previously
- **Grouping** operations
 - Run many operations with one handle
- Adding **assertion** statements
 - Build exceptions into your graph

Grouping




Many variables are updated with separate ops




Running each update op separately is a pain

```
val1, val2, ... = sess.run([update1, update2...])
```




Fix: use dummy op that **depends on all updates**



This gives us simpler (and more semantic) code

```
_ = sess.run(update_all)
```



Grouping Operations “raw”

```
updates = [update1, update2...]  
with tf.control_dependencies(updates):  
    update_all = tf.no_op()
```


TensorFlow has a [built-in helper](#) to make this cleaner

tf.group


```
updates = [update1, update2...]  
update_all = tf.group(*updates)
```

- Uses `tf.control_dependencies` under the hood
- Has extra built-in functionality
 - Automatically **groups operations by device** (CPU, GPU1, GPU2, etc)

Assertions



Need to validate tensor going into this op




Pain in the ass version:

```
_, check_me = sess.run([train, check_op])  
if not validate(check_me):  
    raise ValueError(...)
```


Gets worse the more checks you need to make

Better: **validate as the graph runs**

`tf.assert` **raises exception** if check value is `False`



Make it a dependency to run before critical op



Simple Assertion

```
check_me = tf.multiply(...)
assert_op = tf.assert(check_me != 0, check_me)
with tf.control_dependencies([assert_op]):
    next_op = tf.divide(10, check_me)
```

Required arguments for tf.assert:

1. Boolean check value
2. A tensor to print in the error message


Common scenario: check for NaN or Inf

```
check_me = tf.matmul(...)  
assert_op = tf.check_numerics(check_me, 'It broke!')  
with tf.control_dependencies([assert_op]):  
    next_op = ...
```


Many **built-in assertion helpers**:

1. https://www.tensorflow.org/api_guides/python/check_ops
2. https://www.tensorflow.org/api_guides/python/control_flow_ops#Debugging_Operations

One last example: synchronization



Want to wait for *both* to finish before moving on




tf.tuple

```
wait_1 = tf.some_op(...)
wait_2 = tf.another_op(...)
sync_1, sync_2 = tf.tuple([wait_1, wait_2])
```


Note: TensorFlow already waits for dependencies

`tf.tuple` is generally reserved for unique requirements




CONDITIONAL LOGIC

Idea: different ops based on intermediate results



Like our example from the intro



■ TensorFlow offers two Ops for conditionals

- `tf.cond`
 - Like an `if/else` statement
- `tf.case`
 - Like a `case` statement

Using tf.cond

```
tf.cond(pred, run_if_true, run_if_false)
```

tf.cond takes **three required arguments**

Using tf.cond

```
tf.cond(pred, run_if_true, run_if_false)
```

- The first is a **scalar boolean** predicate
- Switch telling TensorFlow which branch to run

Using tf.cond

```
tf.cond(pred, run_if_true, run_if_false)
```

- The second is a **callable** (function, lambda, etc)
- Should take **no input**, and **return zero or more tensors**
- Runs if the predicate is true

Using tf.cond

```
tf.cond(pred, run_if_true, run_if_false)
```

- The last is also a **callable**
- Similar to previous input
- Runs if the predicate is **false**

Example

```
pred = a < b
def run_if_true():
    return tf.add(3, 3)
def run_if_false():
    return tf.square(3)
out = tf.cond(pred, run_if_true, run_if_false)
```

Define predicate

```
pred = a < b
def run_if_true():
    return tf.add(3, 3)
def run_if_false():
    return tf.square(3)
out = tf.cond(pred, run_if_true, run_if_false)
```

“True” callable

```
pred = a < b
def run_if_true():
    return tf.add(3, 3)
def run_if_false():
    return tf.square(3)
out = tf.cond(pred, run_if_true, run_if_false)
```

“False” callable

```
pred = a < b
def run_if_true():
    return tf.add(3, 3)
def run_if_false():
    return tf.square(3)
out = tf.cond(pred, run_if_true, run_if_false)
```


Put it all together

```
pred = a < b
def run_if_true():
    return tf.add(3, 3)
def run_if_false():
    return tf.square(3)
out = tf.cond(pred, run_if_true, run_if_false)
```


You might one-liner it for simple uses

```
tf.cond(a < b, lambda: tf.add(3, 3), lambda: tf.sqaure(3))
```



The graph we just defined looks like this



The less-than operation outputs a boolean




tf.cond selects the right output to pass along




■ tf.cond notes:

- Both callables' return signatures *must* match
 - Same number of tensors with the same type
- External ops needed for either branch *always* run
 - Place as many ops inside the callables as possible

Now need choose from more than two actions



Same graph, case statement syntax



Instead of chaining `tf.cond` over and over,
we can **use a single `tf.case`**

Using tf.case

```
tf.case(pred_fn_pairs, default)
```

tf.case takes two required arguments

Using tf.case

```
tf.case(pred_fn_pairs, default)
```

- The first is a **list of tuple pairs** (predicate, callable)
- Maps boolean predicates to potential operations to run
- It can also be a dictionary: {pred: callable}

Using `tf.case`

```
tf.case(pred_fn_pairs, default)
```

- The second is a **callable**, as we've seen before
- Runs if none of the predicates are true

Basic example:

```
a = (prev < 0, lambda: prev + 3)
b = (prev < 10, lambda: prev * 3)
c = (prev < 20, lambda: prev - 3)
pairs = [a, b, c]
default = lambda: prev / 3
out = tf.case(pairs, default)
```

Define tuple pairs of predicates/callables

```
a = (prev < 0, lambda: prev + 3)
b = (prev < 10, lambda: prev * 3)
c = (prev < 20, lambda: prev - 3)
pairs = [a, b, c]
default = lambda: prev / 3
out = tf.case(pairs, default)
```

Define a default callable

```
a = (prev < 0, lambda: prev + 3)
b = (prev < 10, lambda: prev * 3)
c = (prev < 20, lambda: prev - 3)
pairs = [a, b, c]
default = lambda: prev / 3
out = tf.case(pairs, default)
```

Create the op


```
a = (prev < 0, lambda: prev + 3)
b = (prev < 10, lambda: prev * 3)
c = (prev < 20, lambda: prev - 3)
pairs = [a, b, c]
default = lambda: prev / 3
out = tf.case(pairs, default)
```

tf.cond notes:

- All callables' return signatures *must* match (like `tf.cond`)
 - Same **number of tensors** with the **same type**
- Only one callable will run
 - As if each **case** has a **break** statement
- Can also pass in attribute **exclusive** (defaults to `False`)
 - Makes op throw exception if more than one predicate is true

■ General notes on conditional logic:

- Ops on non-selected branches **are not run**
 - Great if heavy computation only needs to happen sometimes
 - Example: [stochastic depth](#)
- TensorFlow **differentiates through the selected path**
- For TensorBoard: cond and case can get ugly
 - Use **tf.name_scope** or **tf.variable_scope** inside callables



WHILE LOOPS

Common uses of loops in TensorFlow

1. Feeding intermediate results **back into graph**
2. **“Unrolling”** a loop of operations


Feeding results back into graph

```
my_op = tf.some_op(prev)
...
res = start_val
for i in range(...):
    feed_dict = {prev: res}
    res = sess.run(my_op, feed_dict)
```


Feeding results back into graph

```
my_op = tf.some_op(prev)
...
res = start_val
for i in range(...):
    feed_dict = {prev: res}
    res = sess.run(my_op, feed_dict)
```

The graph looks like this



This loop occurs in the **Python layer**




“Unrolling” loops of operations


```
my_op = tf.some_op(prev)
for i in range(...):
    my_op = tf.some_op(my_op)
```

Basically, create a bunch of ops in the graph


The unrolled graph looks like this




Each additional op **adds overhead**



Ideally: loop in **C++ layer** with **minimal added ops**



`tf.while_loop` is what we're looking for!



tf.while_loop

```
tf.while_loop(cond, body, loop_vars)
```

tf.while_loop takes three required arguments

tf.while_loop

```
tf.while_loop(cond, body, loop_vars)
```

- Let's start with the last: `loop_vars`
- List/tuple of tensors used *in the first iteration* of the while loop
 - The documentation doesn't make this super clear
- These are passed to `both` the condition and body (up next)

tf.while_loop

```
tf.while_loop(cond, body, loop_vars)
```


- Callable. Maps from (*loop_vars) → boolean scalar
- If it returns true, the body executes,
- Otherwise, we exit the loop

tf.while_loop


```
tf.while_loop(cond, body, loop_vars)
```

- Callable. Maps from (*loop_vars) → (*next_loop_vars)
- Main computation takes place here
- Also need to increment counter (if using one) here
- Output from this gets sent to next iteration cond and body


Here is what the basic loop looks like




We pass in our initial loop arguments




Those are now the loop variables




The loop variables get sent to the cond function




If cond is true, we pass the vars to body




body's outputs become the new loop vars




The loop continues while cond evaluates to true



Once cond is false, we return the current loop vars



while_loop compatible with auto-differentiation



Basic while_loop example: 100 loops

```
def cond(i, val):  
    return i < 100  
def body(i, val):  
    return i+1, val + 5  
loop = tf.while_loop(cond, body, (0, 0))
```

Define our condition

```
def cond(i, val):  
    return i < 100  
def body(i, val):  
    return i+1, val + 5  
loop = tf.while_loop(cond, body, (0, 0))
```

Define the body

```
def cond(i, val):  
    return i < 100  
def body(i, val):  
    return i+1, val + 5  
loop = tf.while_loop(cond, body, (0, 0))
```

Build the loop!

```
def cond(i, val):  
    return i < 100  
def body(i, val):  
    return i+1, val + 5  
loop = tf.while_loop(cond, body, (0, 0))
```

Notice that cond and body have same inputs

```
def cond(i, val):  
    return i < 100  
def body(i, val):  
    return i+1, val + 5  
loop = tf.while_loop(cond, body, (0, 0))
```

And that the values are modified in the body

```
def cond(i, val):  
    return i < 100  
def body(i, val):  
    return i+1, val + 5  
loop = tf.while_loop(cond, body, (0, 0))
```

Reusing variables is simple

```
def body(i, val):  
    w = tf.get_variable('w', ...)  
    return i+1, tf.matmul(val, w)
```

Don't have to declare `scope.reuse_variables()`

```
def body(i, val):  
    w = tf.get_variable('w', ...)  
    return i+1, tf.matmul(val, w)
```

Reusing variables + feeding data into itself → **RNN!**

tf.dynamic_rnn is implemented with a `tf.while_loop`

Full implementation beyond scope of lecture


-But small RNN example is in notebook

Optional parameters

- `shape_invariants` (default: None)
 - Allows you to specify which `loop_vars` can have **variable shape**
- `parallel_iterations` (default: 10)
 - Number of allowed parallel iterations (if possible)
- `swap_memory` (default: False)
 - Allows (or disallows) **GPU-CPU memory swap** (RNN backprop is hungry)
- `back_prop` (default: True)
 - Allows (or disallows) backpropagation for this loop.

tf.while_loop notes

- Faster than refeeding with loops of `sess.run()`
 - Roughly 30% improvement
- *Much* faster than unrolling with many static ops
 - Both in graph creation and in run time
- Like with conditionals, TensorBoard graph can get ugly
 - Use name/variable scope for **cond** and **body**



WRAPPING THIS BABY UP

With native control flow:

Data transfer **overhead is minimized**

With native control flow:

Graph logic is **self-contained**

With native control flow:

Enables use of **differentiation and queues**

THANKS!



GitHub: @samjabrahams

Twitter: @sabraha

Email: sam@samabrahams.com

Presentation template by [SlidesCarnival](#)

Copyright 2017 Sam Abrahams. All Rights Reserved.